




Branching & Merging PCollections

Iñigo San José
Miren Esnaola



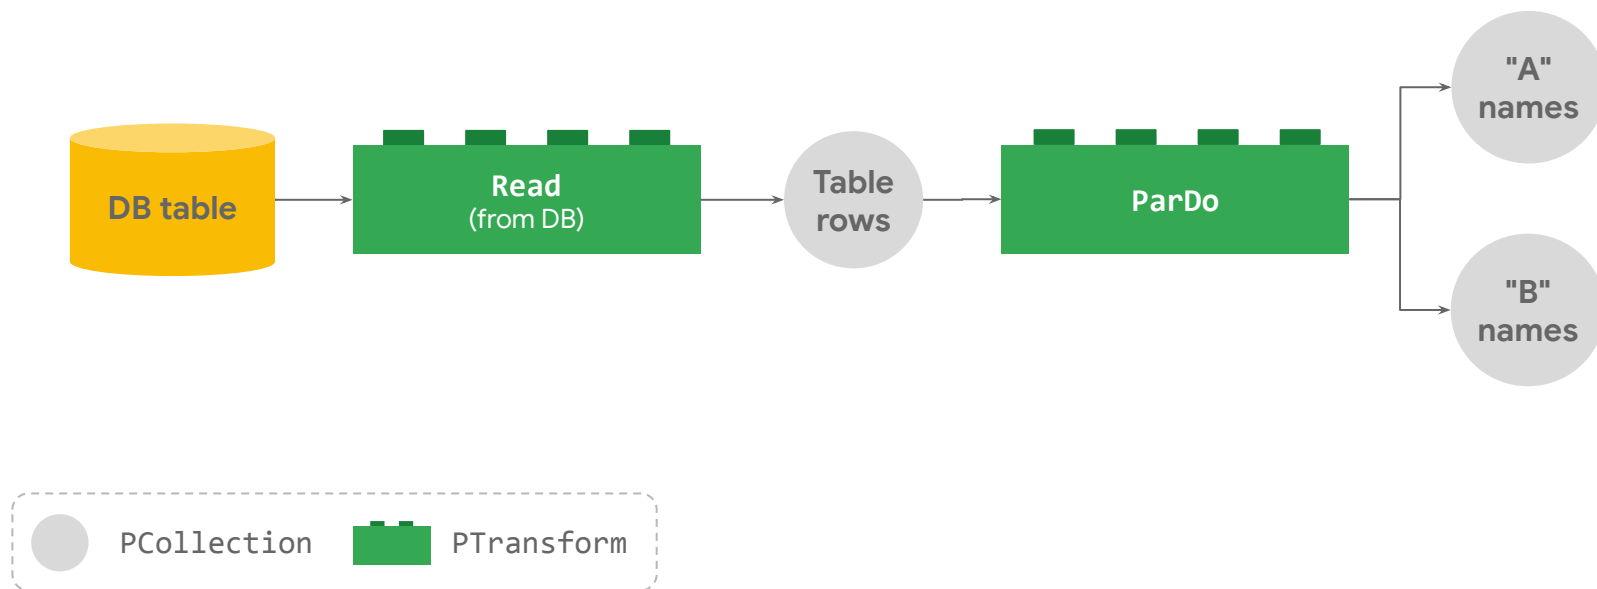
Branching PCollections

Multiple **ParDo** transforms process the same PCollection



Branching PCollections

A single **ParDo** transform producing multiple outputs



Branching PCollections


A **Partition** transform

Partition splits a single collection into a fixed number of smaller collections according to a partitioning function.



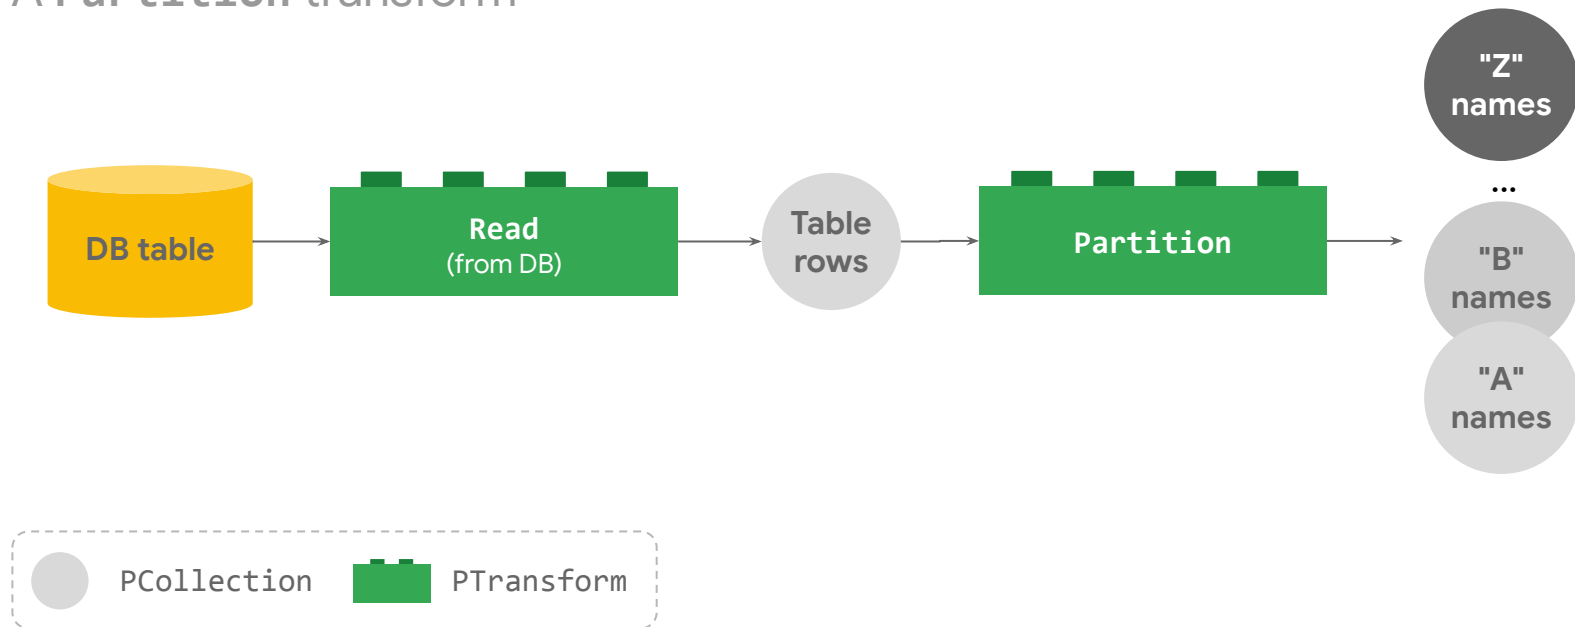
Branching PCollections

A **Partition** transform



Branching PCollections

A **Partition** transform




Merging PCollections

A **Flatten** transform

Flatten is used with collection objects that store the same data type. It merges multiple collection objects into a single logical collection.


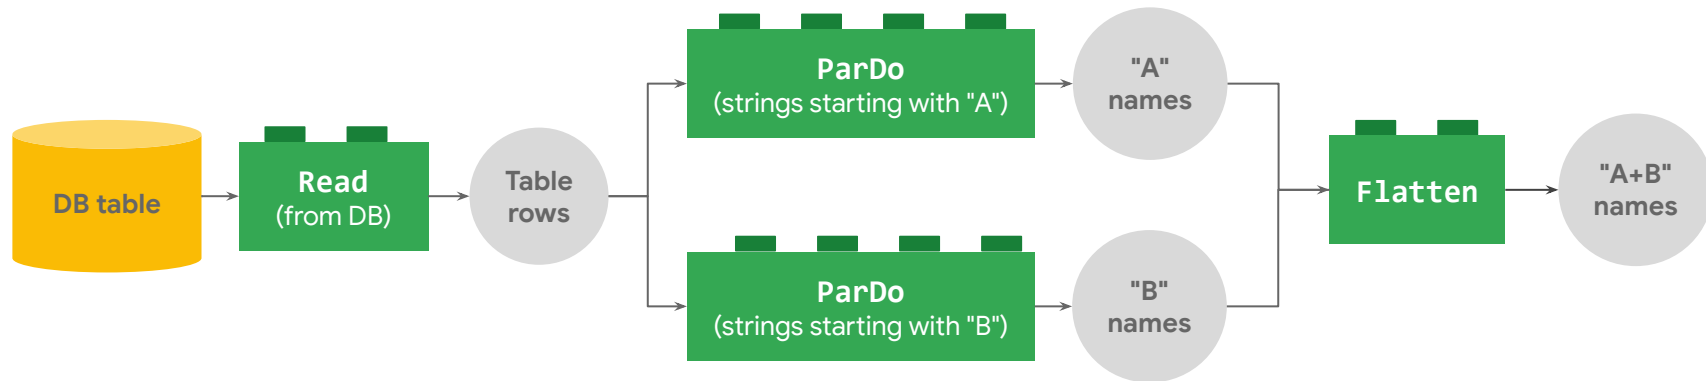
Merging PCollections

A Flatten



Merging PCollections

A Flatten transform



Merging PCollections


A `CoGroupByKey` transform

CoGroupByKey performs a relational join of two or more key/value collections with the same key type.



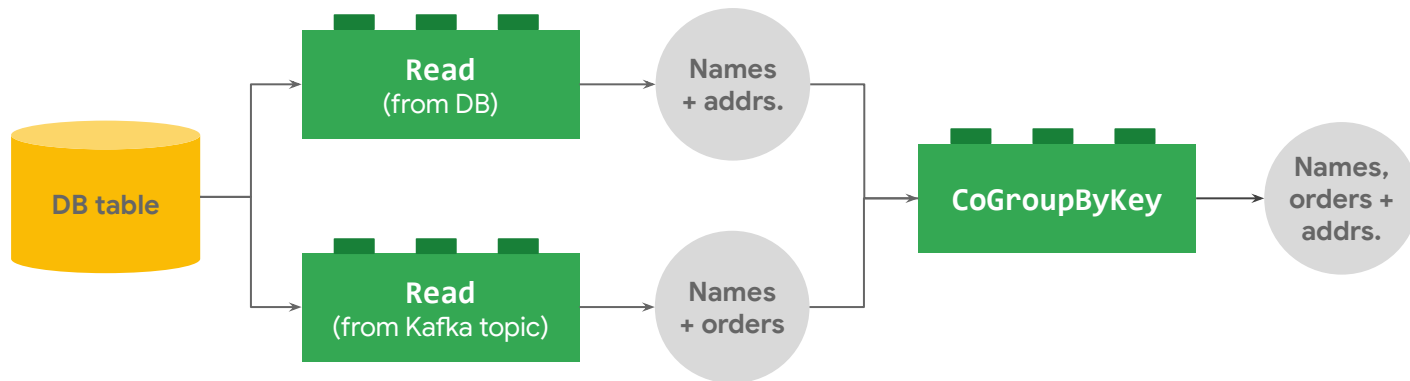
Merging PCollections

A CoGroupByKey transform



Merging PCollections

A CoGroupByKey transform



PCollection



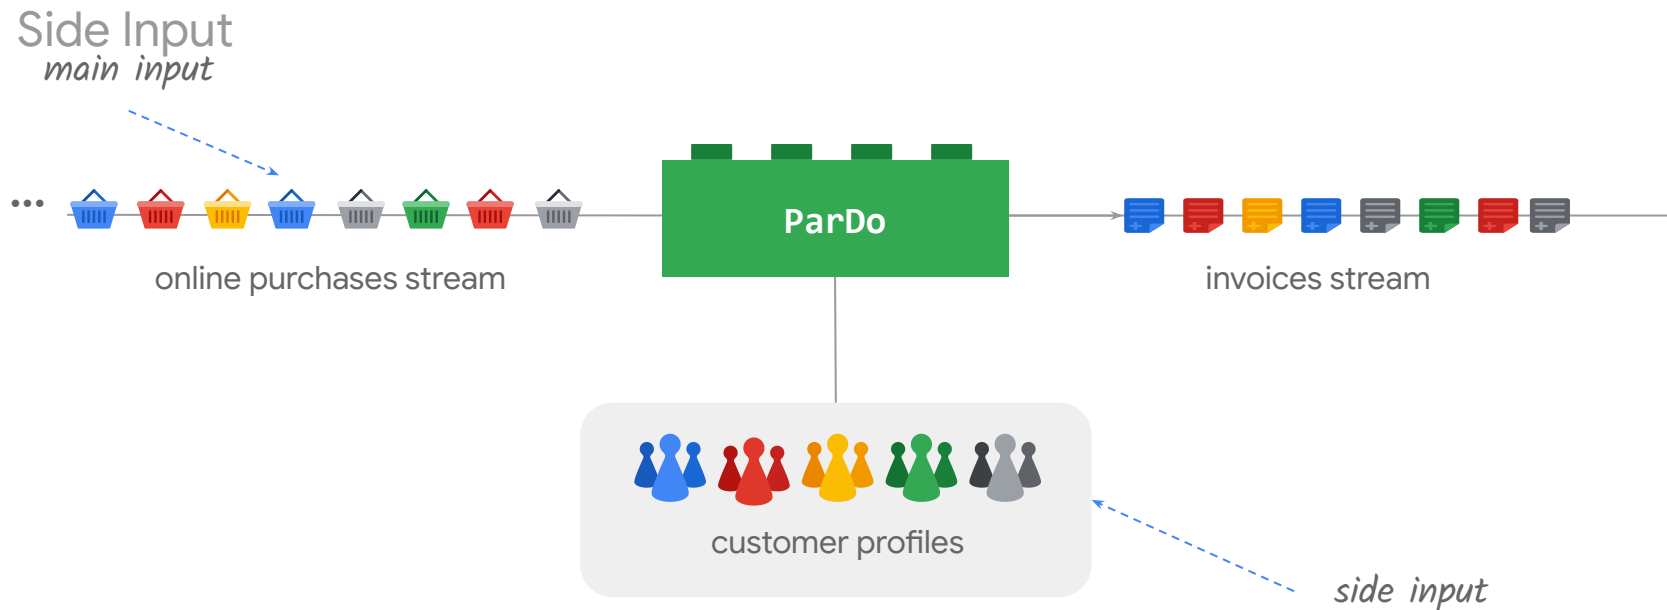
PTransform

Merging PCollections

Side Input

A **side input** is an additional input that a DoFn can access each time it processes an element in the input PCollection. This additional data needs to be determined at runtime (it can not be hard-coded).

Merging PCollections



Side inputs and windowing

- When a view of a windowed `PCollection` is created, the view represents a single entity per window (one singleton per window, one list per window, ...).
- Beam projects the main input element's window into the side input's window set, and then uses the side input from the resulting window.
 - **Identical windows** → projection provides the exact corresponding window.
 - **Different windows** → projection used to choose most suitable side input window.
- If the main input element exists in more than one window, `processElement` gets called once for each window. Each call projects the “current” window for the main input element, and thus might provide a different view of the side input each time.
- If the side input has multiple trigger firings, the value from the latest trigger firing is used.

Side inputs and windowing

As the sales season goes by, the discounts are larger week by week...

Side input

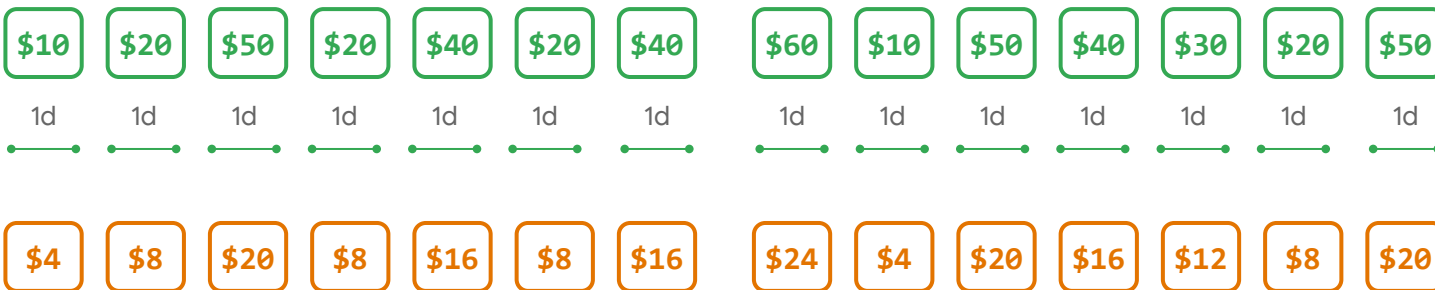
60% discount

70% discount

1w

1w

Main input



Thank you!

Questions?

