



# Dataflow Templates

Akvelon

<https://akvelon.com/>





# Agenda

Dataflow Templates

Use Case - Data Protection Using Tokenization

Solution Architecture & Technical Highlights

Demo

Contributing to open source [DataflowTemplates](#)





Google Cloud  
**Partner**

# Beam Pipeline





# Dataflow Templates

Template - a way to package and stage pipelines

Template types: Classic and Flex

Flex template benefits

- Dynamic DAG
- Eliminated need for ValueProvider
- Expanded templates flexibility

The screenshot shows the Google Cloud Platform interface for creating a Dataflow job from a template. The left sidebar shows the navigation menu with 'Dataflow' selected. The main panel is titled 'Create job from template'. It contains several form fields: 'Job name' (set to 'Demo'), 'Regional endpoint' (set to 'us-central1'), and 'Dataflow template' (set to 'Custom Template'). Below these, there is a section for 'Template path' with a checked checkbox and a path 'gs://tokenization\_test/templates/datatokenization\_test.json'. Further down, there is a 'Required parameters' section with three input fields: 'GCS location of data schema', 'GCS file pattern for input files', and 'File format of input files'. At the bottom, there is a checkbox for 'CSV file(s) contain headers'.



# Create Template

- Implement pipeline
- Create metadata
- Build template

**Job name \***

Must be unique among running jobs

**Regional endpoint \***

us-central1

Choose a Dataflow regional endpoint to deploy worker instances and store job metadata. You can optionally deploy worker instances to any available Google Cloud region or zone by using the worker region or worker zone parameters. Job metadata is always stored in the Dataflow regional endpoint. [Learn more](#)

**Dataflow template \***

Word Count

Batch pipeline. Reads text from Cloud Storage, tokenizes text lines into individual words, and performs frequency count on each of the words. [OPEN TUTORIAL](#)

## Required parameters

**Input file(s) in Cloud Storage \***

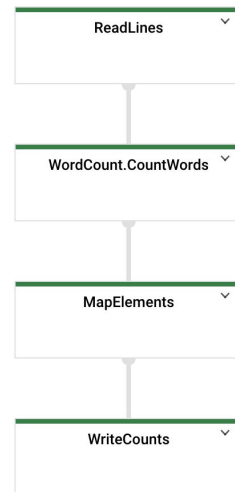
The input file pattern Dataflow reads from. Use the example file (gs://dataflow-samples/shakespeare/kinglear.txt) or enter the path to your own using the same format: gs://your-bucket/your-file.txt

**Output Cloud Storage file prefix \***

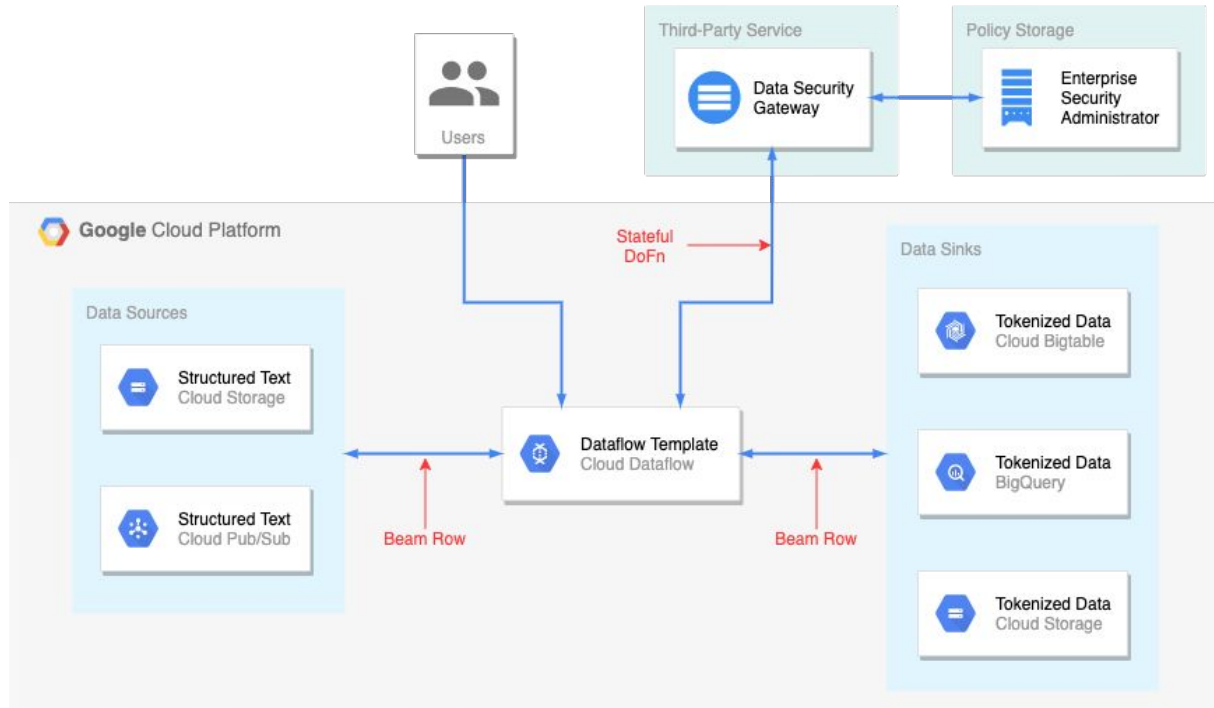
Path and filename prefix for writing output files. Ex: gs://your-bucket/counts

**Temporary location \***

Path and filename prefix for writing temporary files. Ex: gs://your-bucket/temp

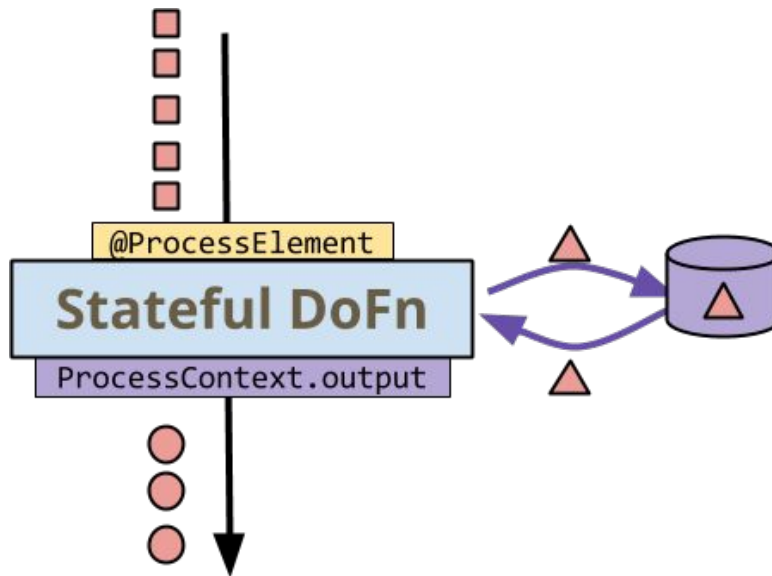


# Data Protection Using Tokenization



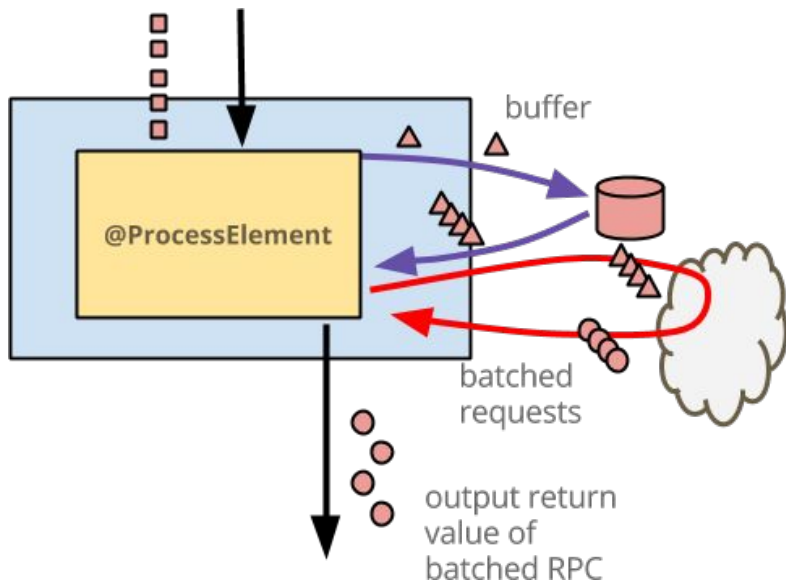


# Stateful Processing



More info <https://beam.apache.org/blog/stateful-processing/>

# Stateful DoFn



```
new DoFn<Event, EnrichedEvent>() {

    private static final int MAX_BUFFER_SIZE = 500;

    @StateId("buffer")
    private final StateSpec<BagState<Event>> bufferedEvents = StateSpecs.bag();

    @StateId("count")
    private final StateSpec<ValueState<Integer>> countState = StateSpecs.value();

    @ProcessElement
    public void process(
        ProcessContext context,
        @StateId("buffer") BagState<Event> bufferState,
        @StateId("count") ValueState<Integer> countState) {

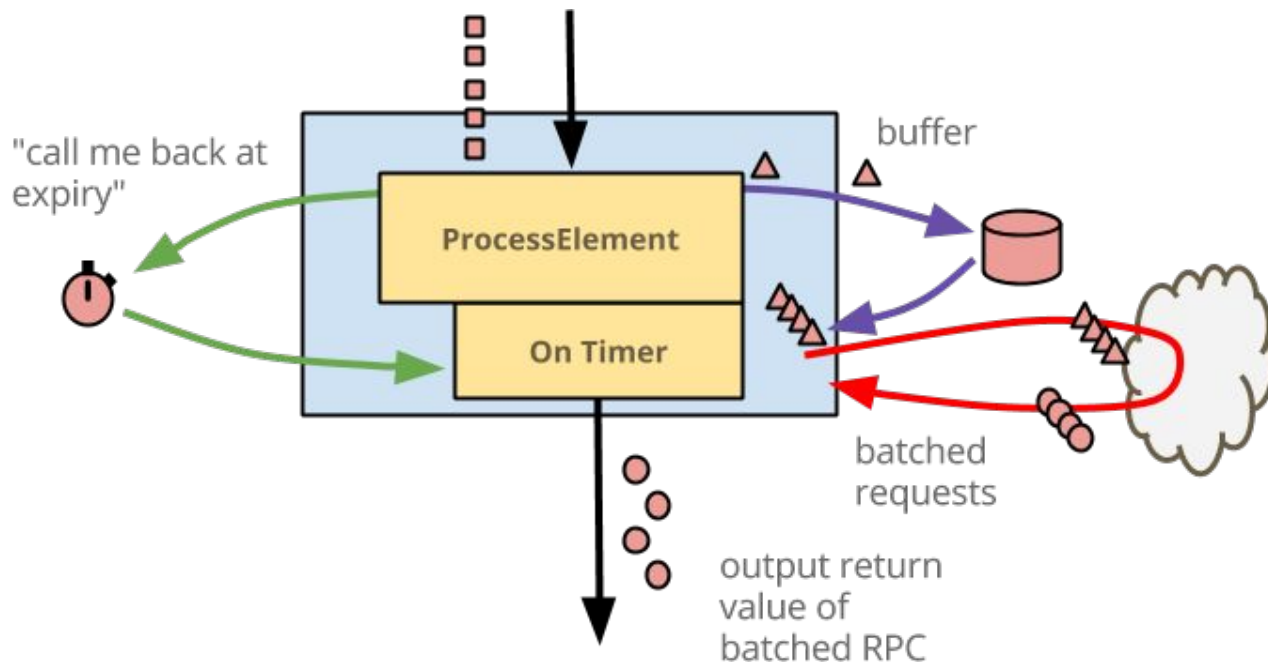
        int count = firstNonNull(countState.read(), 0);
        count = count + 1;
        countState.write(count);
        bufferState.add(context.element());

        if (count >= MAX_BUFFER_SIZE) {
            for (EnrichedEvent enrichedEvent : enrichEvents(bufferState.read())) {
                context.output(enrichedEvent);
            }
            bufferState.clear();
            countState.clear();
        }
    }

    // ... TBD ...
};
```

More info <https://beam.apache.org/blog/timely-processing/>

# Timely Stateful Processing



More info <https://beam.apache.org/blog/timely-processing/>

# Stateful DoFn with Timer - Implementation

```
new DoFn<Event, EnrichedEvent>() {  
    //...  
    @TimerId("expiry")  
    private final TimerSpec expirySpec = TimerSpecs.timer(TimeDomain.EVENT_TIME);  
  
    @ProcessElement  
    public void process(  
        ProcessContext context,  
        BoundedWindow window,  
        @StateId("buffer") BagState<Event> bufferState,  
        @StateId("count") ValueState<Integer> countState,  
        @TimerId("expiry") Timer expiryTimer) {  
        expiryTimer.set(window.maxTimestamp().plus(allowedLateness));  
  
        //... same logic as before ...  
    }  
  
    @OnTimer("expiry")  
    public void onExpiry(  
        OnTimerContext context,  
        @StateId("buffer") BagState<Event> bufferState) {  
        if (!bufferState.isEmpty().read()) {  
            for (EnrichedEvent enrichedEvent : enrichEvents(bufferState.read())) {  
                context.output(enrichedEvent);  
            }  
            bufferState.clear();  
        }  
    }  
};
```



# GroupIntoBatches

- Groups your data into batches
- Implemented using Stateful DoFn
- Has several optimizations
  - Prefetches data
  - Autosharding in Dataflow

```
import apache_beam as beam

with beam.Pipeline() as pipeline:
    batches_with_keys = (
        pipeline
        | 'Create produce' >> beam.Create([
            ('spring', '🍎'),
            ('spring', '🥕'),
            ('spring', '🍆'),
            ('spring', '🍅'),
            ('summer', '🥕'),
            ('summer', '🍅'),
            ('summer', '🍆'),
            ('fall', '🥕'),
            ('fall', '🍅'),
            ('winter', '🍆'),
        ])
        | 'Group into batches' >> beam.GroupIntoBatches(3)
        | beam.Map(print))
```

Output:

```
('spring', ['🍎', '🥕', '🍆'])
('summer', ['🥕', '🍅', '🍆'])
('spring', ['🍅'])
('fall', ['🥕', '🍅'])
('winter', ['🍆'])
```

More info

<https://beam.apache.org/documentation/transforms/python/aggregation/groupintobatches>



# Processing with GroupIntoBatches

```
public PCollectionTuple expand(PCollection<KV<Integer, Row>> inputRows) {
    FailsafeElementCoder<Row, Row> coder =
        FailsafeElementCoder.of(RowCoder.of(schema()), RowCoder.of(schema()));

    Duration maxBuffering = Duration.millis(MAX_BUFFERING);
    PCollectionTuple pCollectionTuple =
        inputRows
            .apply(
                name: "GroupRowsIntoBatches",
                GroupIntoBatches.<Integer, Row>ofSize(batchSize())
                    .withMaxBufferingDuration(maxBuffering))
            .apply(
                name: "Tokenize",
                ParDo.of(new TokenizationFn(schema()), rpcURI(), failureTag())
                    .withOutputTags(successTag(), TupleTagList.of(failureTag())));

    return PCollectionTuple.of(
        successTag(), pCollectionTuple.get(successTag()).setRowSchema(schema())
        .and(failureTag(), pCollectionTuple.get(failureTag()).setCoder(coder));
}
```



# Use Case Recommendations

## Stateful DoFn

Customization of stateful processing

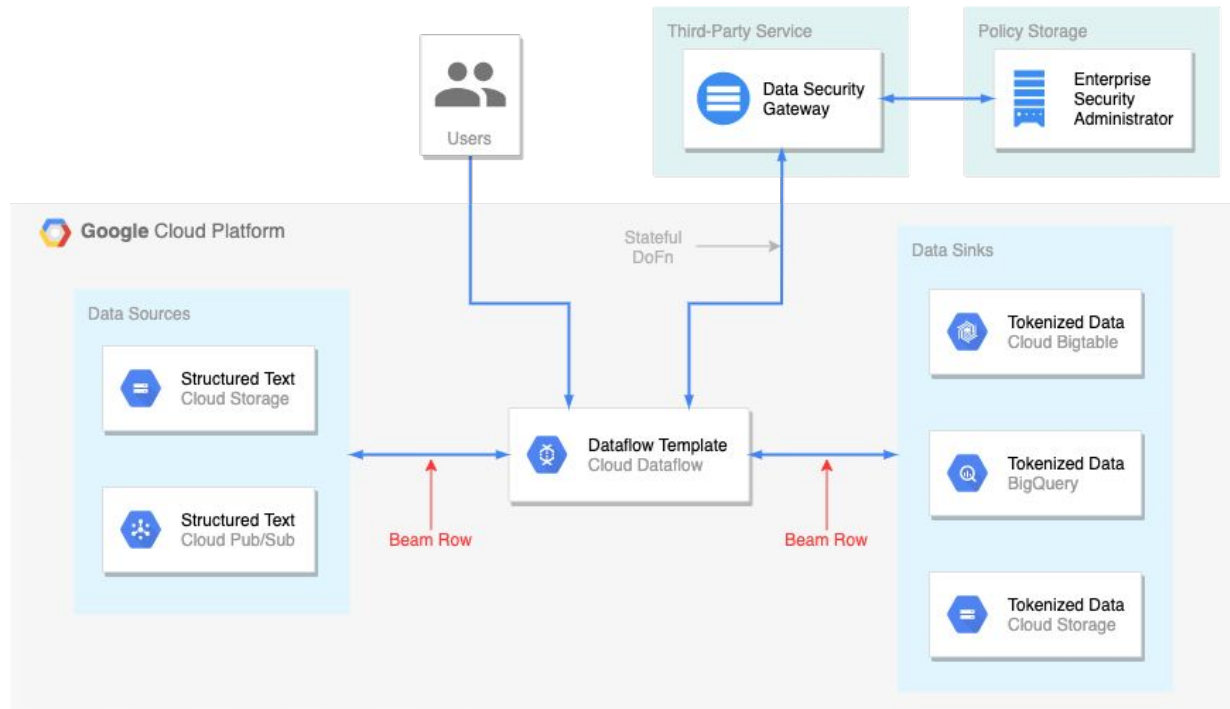
- Deduplication
- Arbitrary-but-consistent indexing

## GroupIntoBatches

Optimized out-of-the-box transform

- External API call
- Data preparation for ML model

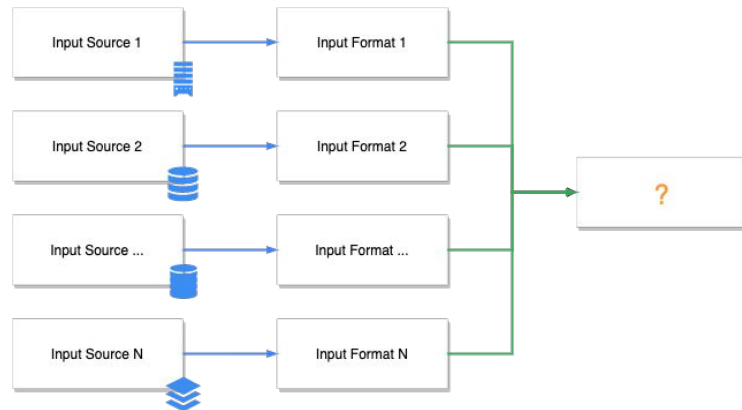
# Data Protection Using Tokenization





# How to Represent Data in Pipeline Code?

- Common abstractions for data representation
- Avoid writing transformation again and again
- Common interface to all IO transforms
- Easy and effective to serialize



# Common Approaches

## Text based formats

- XML
- JSON
- CSV
- YAML

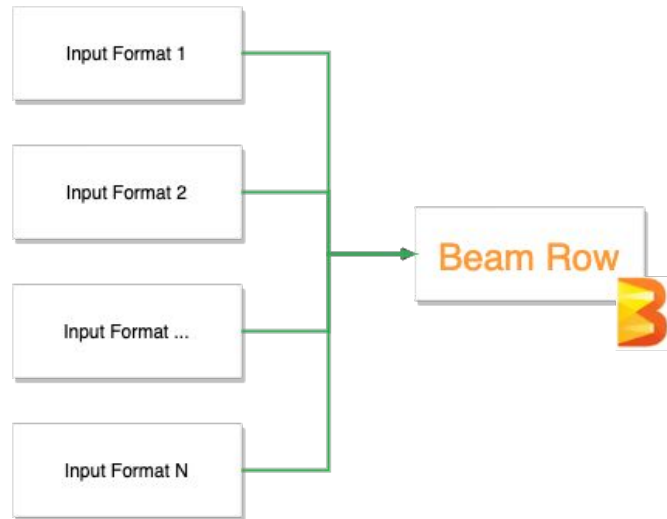
## Binary based formats

- Apache Avro
- Protobuf
- MessagePack

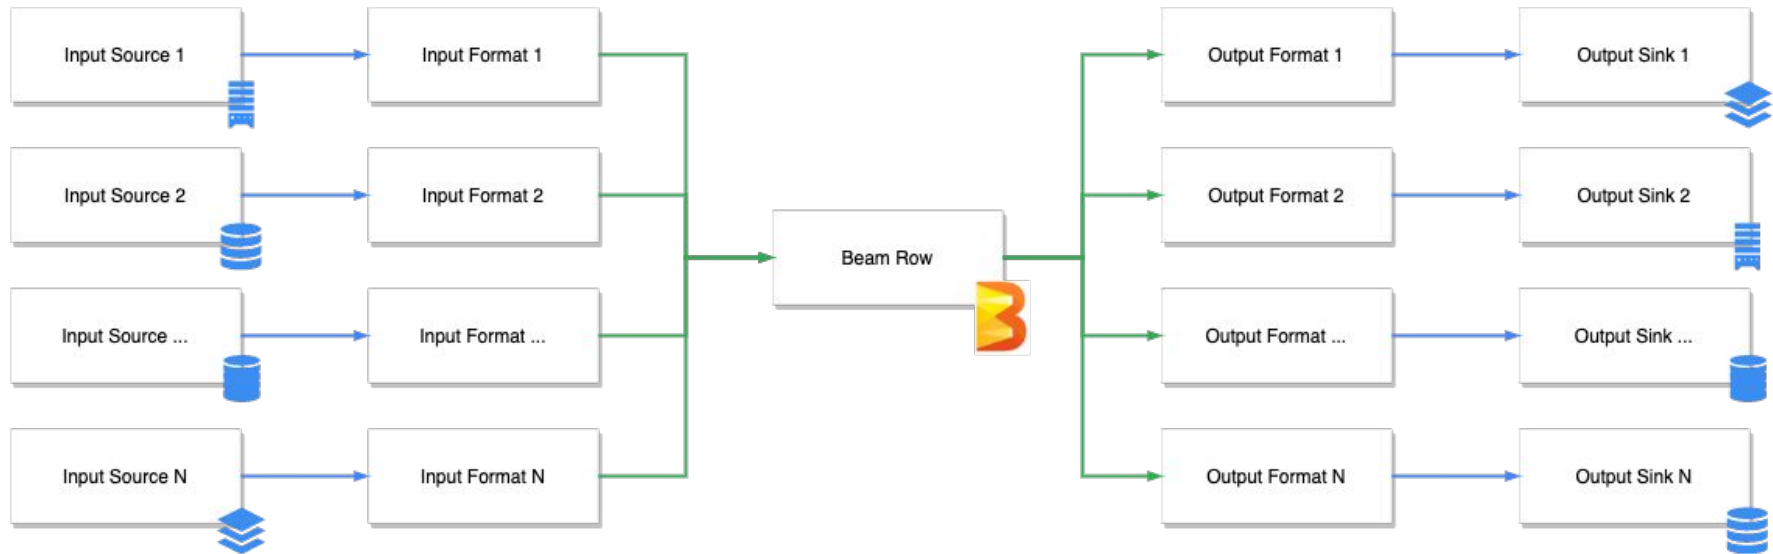
**BSON** **JSON**  
**YAML** **ORC** **PROTOBUF**  
**AVRO** **PARQUET** **XML**  
**CSV** **THRIFT**  
**MESSAGEPACK**

# Beam Row!

- A schema which supports primitives
- Ordering same with schema
- Quite effective serialization using RowCoder
- Generic for json specified schema
- Available for Beam SQL



# How We Use Beam Row?





# Demo

- Template overview
- Add a new input source Parquet IO
- Implement Transform Parquet to Row
- Build and Run template



# Demo

1

—

2

—

3

—

4

## Template Overview



Java Template structure



Metadata file



Schemas



IO transforms

## Add a New Input Source Parquet IO



Read .parquet files



Transform from Parquet to Beam Row



Work with schemas

## Build and Run Template



Template building



Ways to run template

Visit codelab at <https://github.com/akvelon/codelabs>



# Contributing to Dataflow Templates

[https://github.com/GoogleCloudPlatform/  
DataflowTemplates](https://github.com/GoogleCloudPlatform/DataflowTemplates)

1. Fork the repository to develop your template
2. Follow [style guides](#) and [best practices](#)
3. Sign CLA
4. Create a PR
5. LGTM code review!

The screenshot shows the GitHub repository for GoogleCloudPlatform/DataflowTemplates. The repository is named "DataflowTemplates" and is located under the "GoogleCloudPlatform" organization. It has 62 stars, 590 forks, and 415 issues. The repository is currently on the "master" branch. The file explorer shows the following files: "src", "README.md", and "pom.xml". The "README.md" file is selected, and its content is displayed below. The README describes a Dataflow Flex Template for ingesting data from Apache Kafka to Google Cloud Pub/Sub. It lists supported data formats (JSON, PubSubMessage), supported input source configurations (Kafka bootstrap servers, SASL/SCRAM authentication, HashiCorp Vault), supported destination configurations (Google Pub/Sub topic), and supported SSL certificate locations (Google Cloud Storage Bucket). It also provides a brief overview of the template's functionality and the authentication requirements for different scenarios.

GoogleCloudPlatform / DataflowTemplates

<> Code 74 Issues 12 Pull requests 2 Actions Projects 2 Security Insights

master DataflowTemplates / v2 / kafka-to-pubsub / Go to file Add file ...

prathapreddy123 and cloud-teleport PR #176: Kafka to PubSub template Feb 198 on Dec 3, 2020 History

File	Commit	Time
src	PR #176: Kafka to PubSub template	3 months ago
README.md	PR #176: Kafka to PubSub template	3 months ago
pom.xml	PR #176: Kafka to PubSub template	3 months ago

README.md

## Dataflow Flex Template to ingest data from Apache Kafka to Google Cloud Pub/Sub

This directory contains a Dataflow Flex Template that creates a pipeline to read data from a single or multiple topics from [Apache Kafka](#) and write data into a single topic in [Google Pub/Sub](#).

Supported data formats:

- Serializable plaintext formats, such as JSON
- [PubSubMessage](#)

Supported input source configurations:

- Single or multiple Apache Kafka bootstrap servers
- Apache Kafka SASL/SCRAM authentication over plaintext or SSL connection
- Secrets vault service [HashiCorp Vault](#)

Supported destination configuration:

- Single Google Pub/Sub topic

Supported SSL certificate location:

- Bucket in [Google Cloud Storage](#)

In a simple scenario, the template will create an Apache Beam pipeline that will read messages from a source Kafka server with a source topic, and stream the text messages into specified Pub/Sub destination topic. Other scenarios may need Kafka SASL/SCRAM authentication, that can be performed over plain text or SSL encrypted connection. The template supports using a single Kafka user account to authenticate in the provided source Kafka servers and topics. To support SASL authentication over SSL, the template will need access to a secrets vault service with Kafka username and password, and with SSL certificate location in Google Cloud Storage Bucket, currently supporting HashiCorp Vault.



# Summary

- Dataflow Flex templates package pipelines into containers
- Stateful processing in Apache Beam
- BeamRow provides flexible abstraction for data representation
- [DataflowTemplates](#) repository - Google and community contributed templates and utilities





# Thank You!



<https://akvelon.com/>

[@AkvelonInc](#)

