Dataflow Logging

Mehran Nazir
Product Manager, Google Cloud
https://www.linkedin.com/in/mehrannazir/

Outline

Introduce Dataflow's logging & error reporting integration -- we make it easy to bring the best of Cloud Logging with your Dataflow application

Find the logging panel at the bottom
- Step logs vs. worker logs
- Severity types of logs
- Show them Cloud Logging interface
- Explain indicators for specific log name

Diagnostics tabs
- Show Error Reporting interface
- Show them links to platform errors

Advanced things you can do with logs:
- Custom logs using SLF4J

Discuss costs:
- Dataflow service logs are not chargeable
- Which logs are chargeable?
- How do you control that?

Demo
- Walk them through the Logging Panel
- Show them the step logs vs. worker logs
- Show them filtering severity levels of logs
- Click over to Cloud Logging -- show a sample query for logs
- Click over to Error Reporting -- show them diagnostics for errors

Logging & Error Reporting

Every SRE best practices book will emphasize the importance of having robust logging & error reporting capabilities

Google Cloud offers an end-to-end suite for these needs with Operations Suite

- Cloud Monitoring
- Cloud Logging
- Error Reporting
- Cloud Profiler
- Cloud Trace
- Cloud Debugger
- … and several more

What do I need to do to extend these to our Dataflow applications?

Dataflow + Logging & Error Reporting

Nothing at all!

Dataflow offers native integration with Cloud Logging and Error Reporting

Dataflow Logging

Logging Panel

You'll find the logging panel at the bottom of your Job Details page:

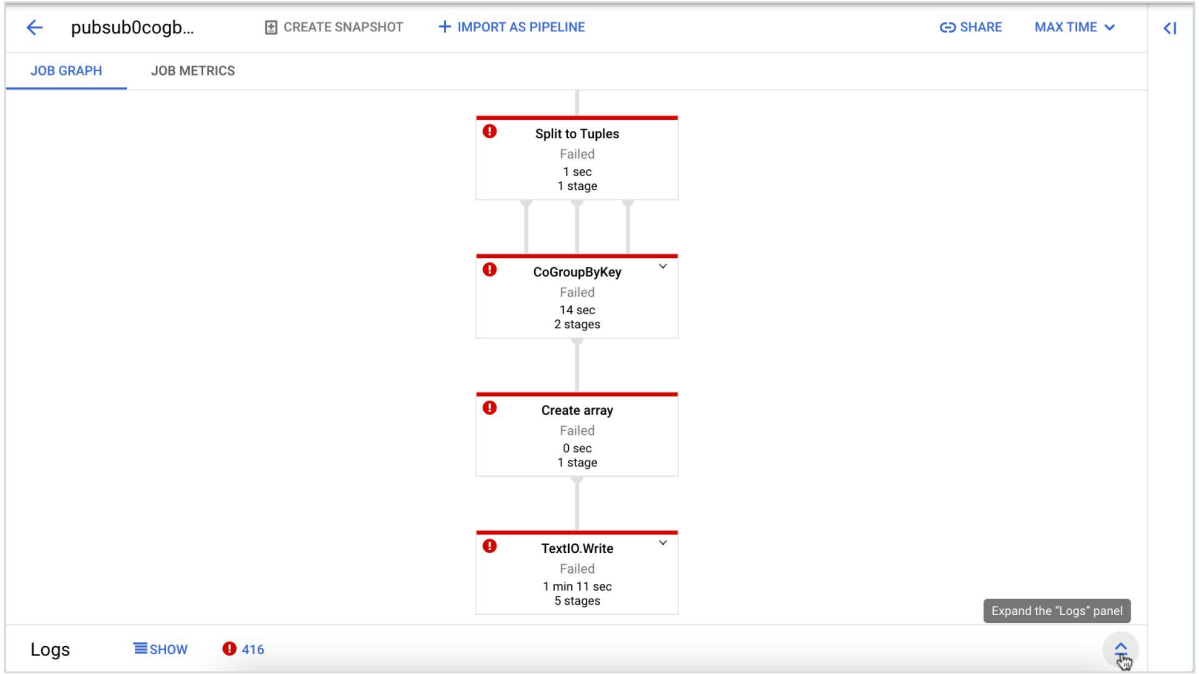

Logging Panel

You'll find three panels:

- Job Logs
- Worker Logs
- Diagnostics

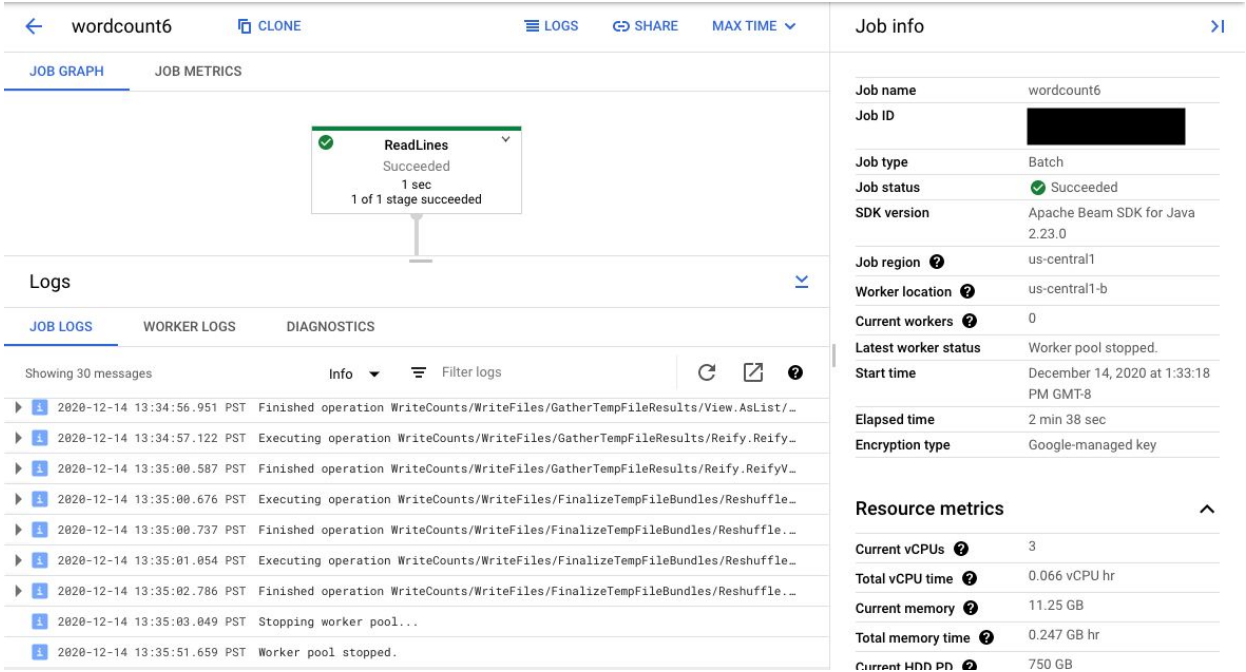

Logs

Job Logs

Messages that report the status of a job as a whole

Worker Logs

Messages that are produced by the worker instances

Diagnostics

Frequency of each error across time observed in your job

Logging Panel - A Closer Look

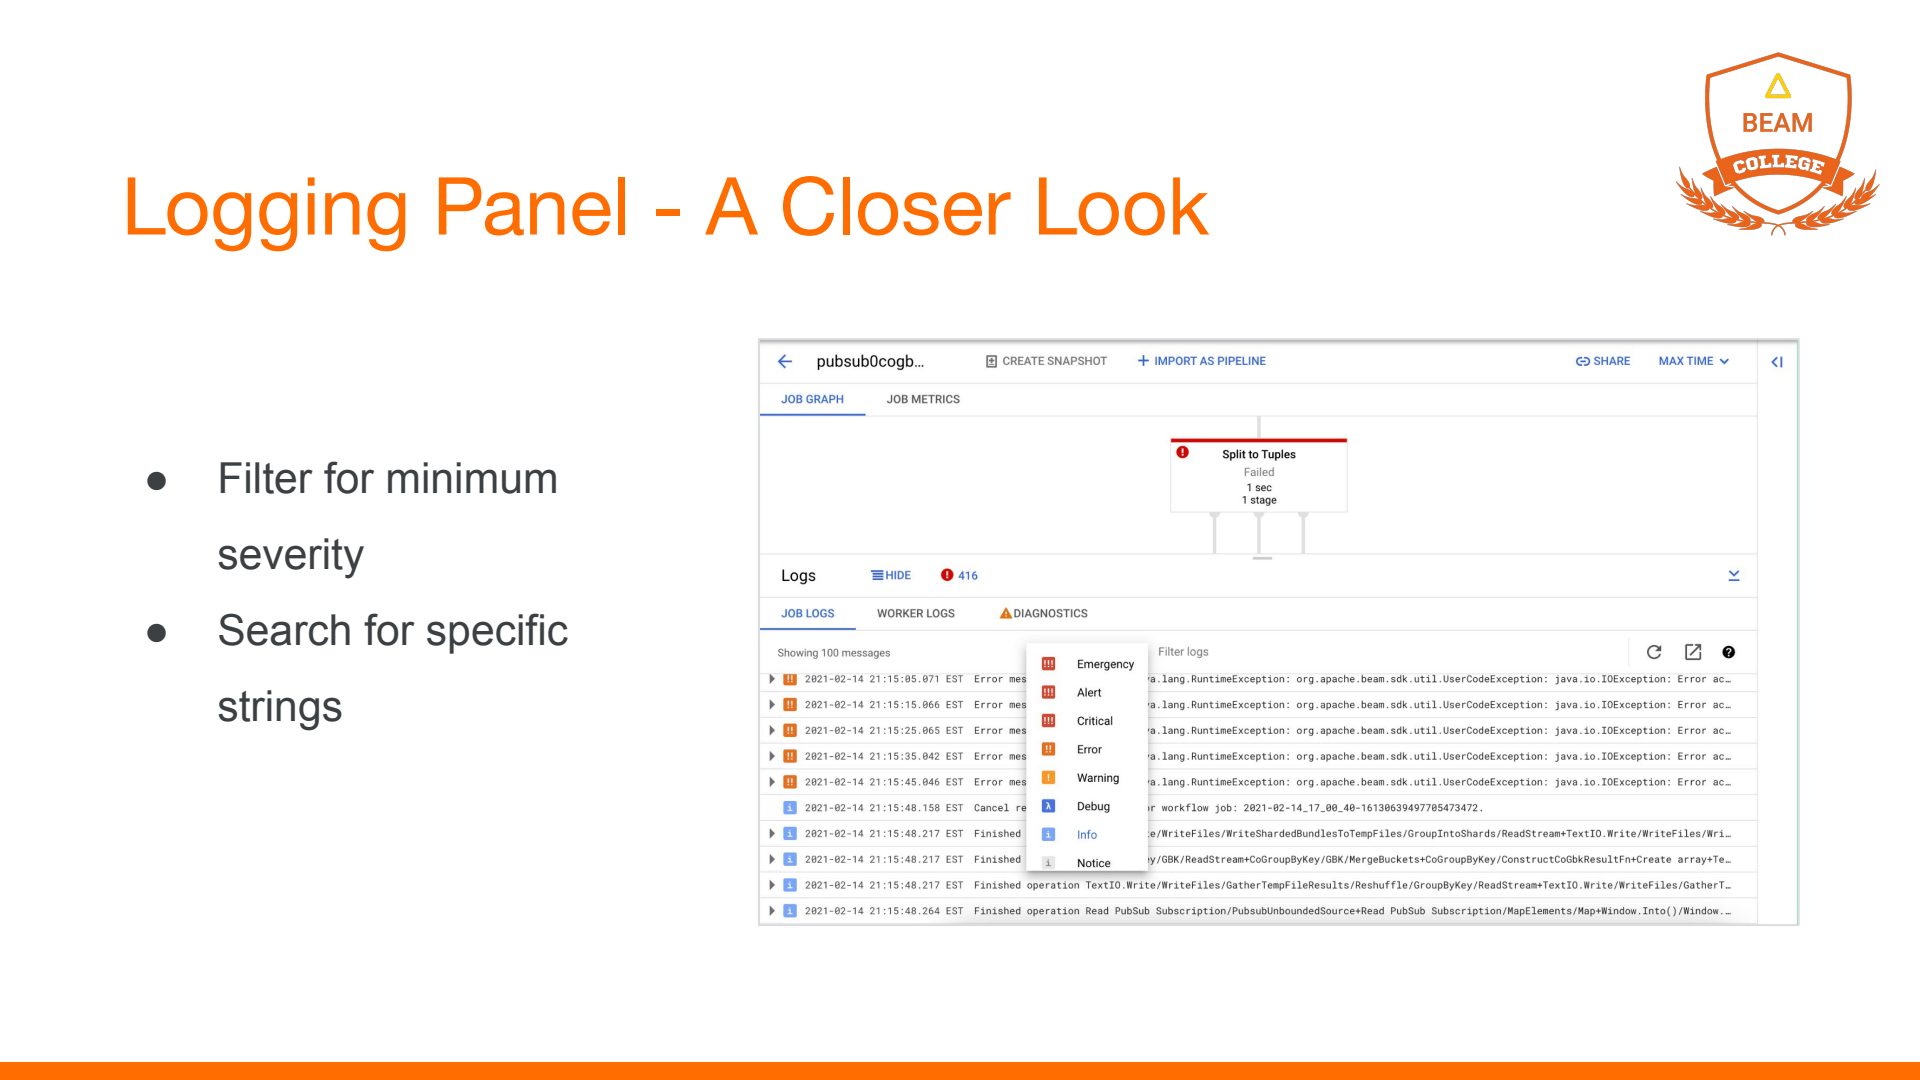

- Filter for minimum severity
- Search for specific strings

Searching Logs

Initialization failure?	`Worker jar file misconfiguration`	
Memory pressure?	`Out of memory: Kill process`	`Shutting down JVM after consecutive periods of measured GC thrashing`
Slow processing?	`Lengthy operation in step`	`Hot key detected`
Large amount of data?	`Commit Key request Exceeds Size Limit`	
Too much logging?	`Throttling logger worker`	

Logs Explorer

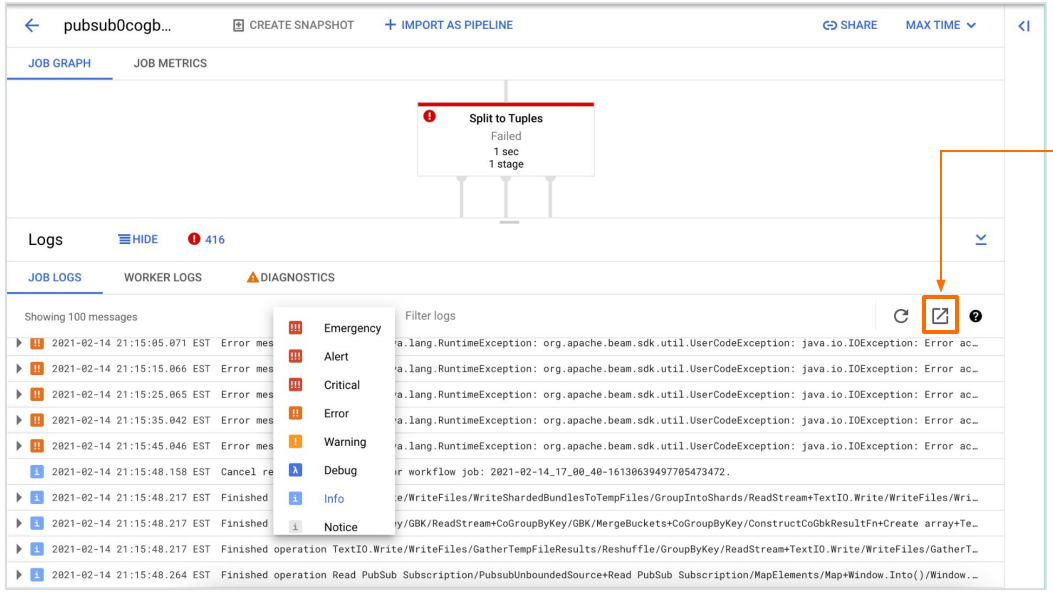

Link to Logs Explorer

Logs Explorer

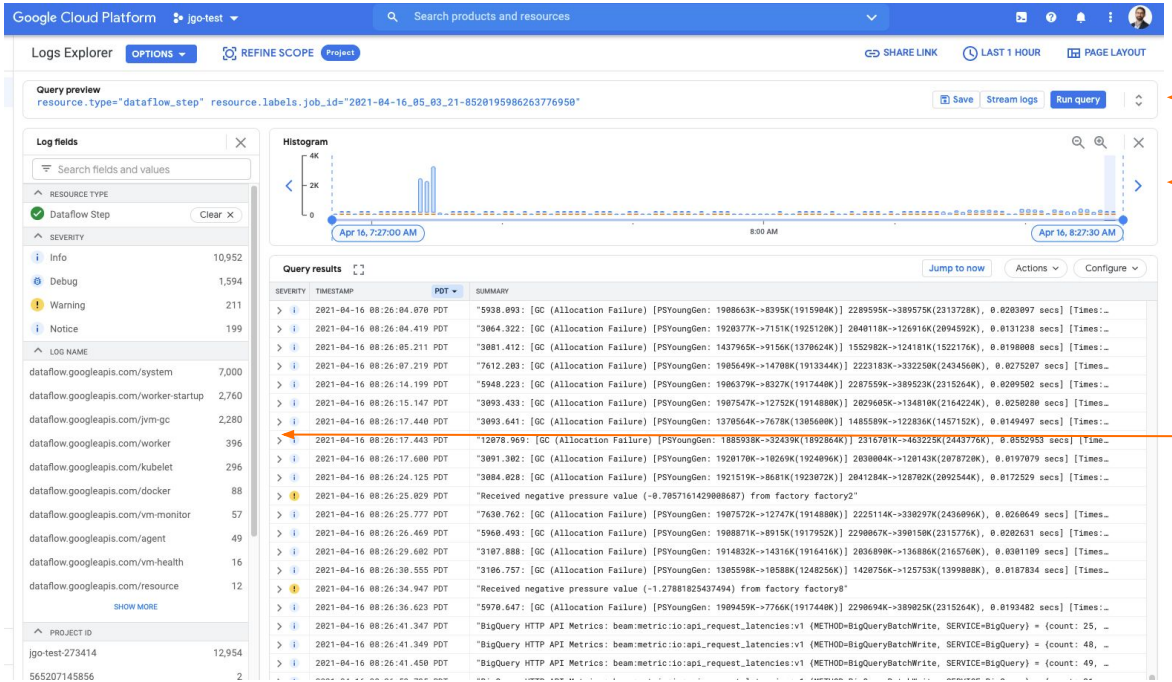

Create custom queries for your logs

Observe the frequency different types of logs

Filter for different types of logs

Different Types of Logs

- **job-message:** job-level messages that various components of Dataflow generate (i.e. autoscaling configuration, when workers start up or shut down, progress on job step, and job errors). Worker-level errors that originate from crashing user cude and that are present in **worker** logs also propagate up to the job-message logs.
- **worker:** messages produced by Dataflow workers. Workers do most of the pipeline work (i.e. they are the only applying ParDos to data). Worker logs contain messages logged by your code and Dataflow.
- **worker-startup:** captures messages related to the startup process. This includes downloading a job's jars from Cloud Storage, then starting the workers. If there is a problem starting workers, these logs are a good place to look.
- **shuffler:** messages from workers that consolidate the results of parallel pipeline operations.
- **docker & kubelet:** messages related to processes by Dataflow service

Custom Logging

You can also (and should!) write custom logs to track events of note

Java

Apache Beam provides the open source SLF4J library

Python

Apache Beam provides the `logging` library package

```
PCollectionTuple eventsTuple = events
                .apply("Filter bad events", ParDo.of(new DoFn<Event, Event>() {

                    private final Counter counter = Metrics.counter(EventProcessingPipeline.class,
"bad-counter");

                    @ProcessElement
                    public void processElement(ProcessContext ctx) {
                        Event event = ctx.element();
                        Double dataValue = event.getDataValue();
                        if (dataValue < 10) {
                            counter.inc();
                            if (dataValue < 0) {
                                LOG.warn("Received negative pressure value ({}) from factory {}",
dataValue, event.getFactoryCode());
                            }
                            ctx.output(badEventsTupleTag, event);
                        }
                        else {
                            ctx.output(goodEventsTupleTag, event);
                        }
                    }
                }).withOutputTags(goodEventsTupleTag, TupleTagList.of(badEventsTupleTag)));
```

Custom Logs in Logs Explorer

- You can find custom logs in the Logs Explorer with:
  - resource.type: dataflow_step
  - logName: dataflow.googleapis.com%2Fworker

Dataflow Error Reporting

Diagnostics Tab

Diagnostics tab shows:

- Where errors occurred along the chosen timeline
- Count of all logged errors
- Possible recommendations for your pipeline

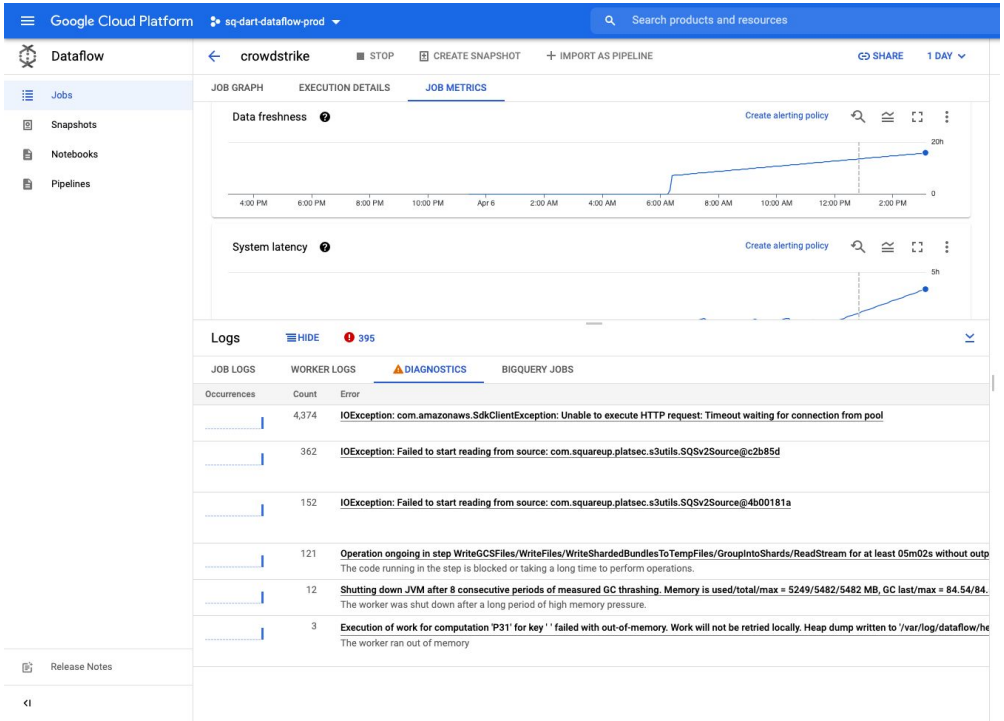

Diagnostics Tab - Recommendations

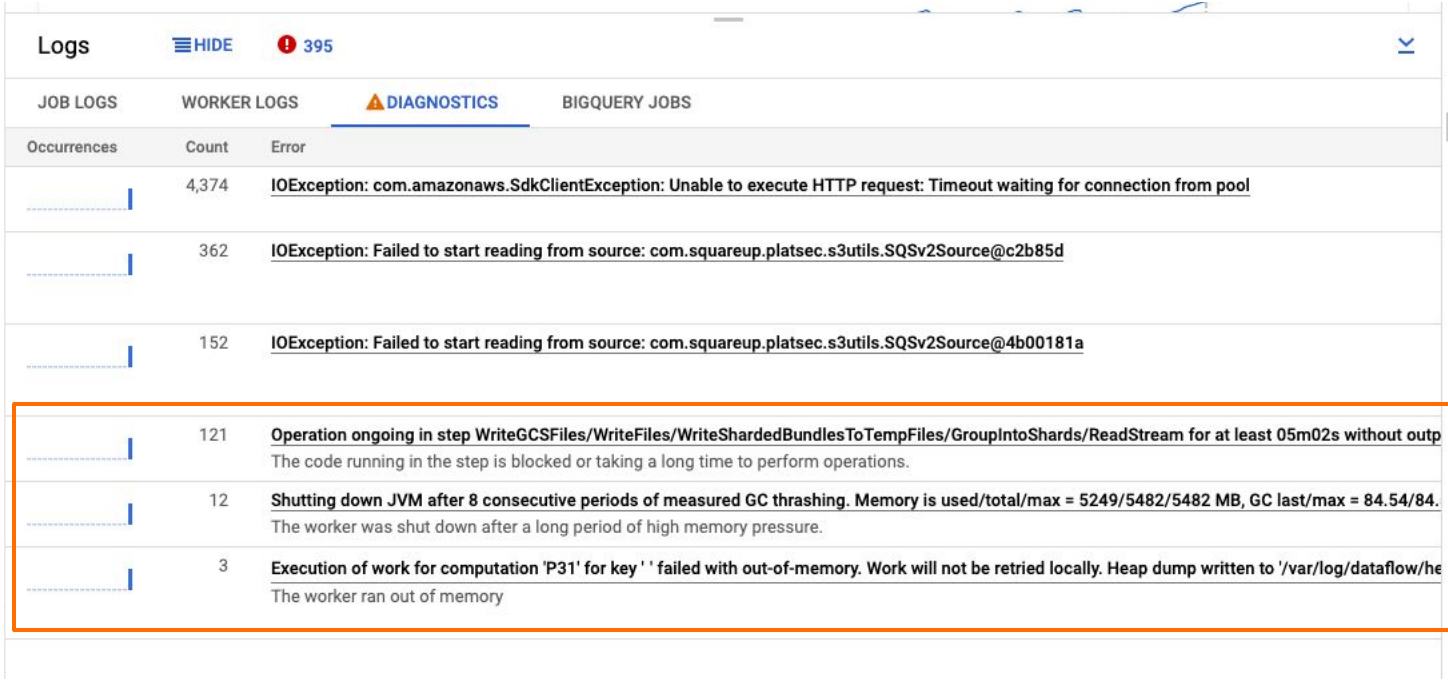

Error Reporting Interface

- Histogram showing errors across your Dataflow project
- Different time ranges
- Look at errors across jobs
- Examine stack traces
- Link to internal issue tracker

Best Practices

Dead-letter queues & Error logging

- Write erroneous records to a dead-letter queue
- Use an exception block that allows you to log an issue & send the raw data as a SideOutput to store unprocessable data

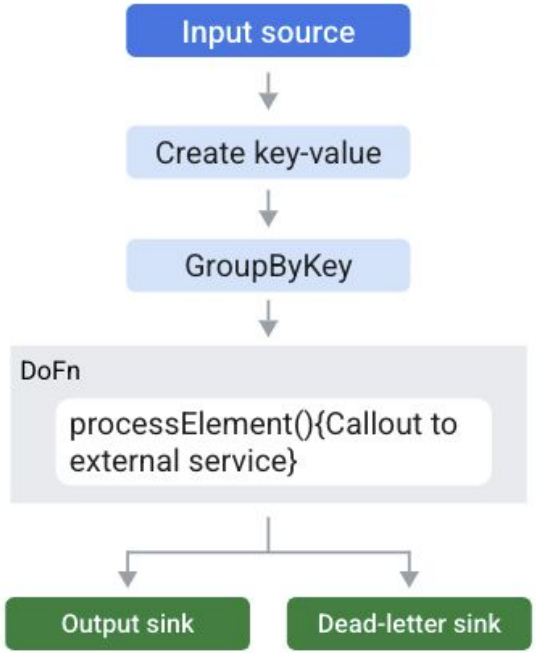

```
final TupleTag successTag;
final TupleTag deadLetterTag;
PCollection input = /* ... */;

PCollectionTuple outputTuple = input.apply(ParDo.of(new DoFn(){
    @Override
    void processElement(ProcessContext ctxt) {
        try {
            c.output(process(c.element));
        } catch(MyException ex) {
            // Optional Logging at debug level
            c.sideOutPut(deadLetterTag, c.element);
        }
    }
})).writeOutPutTags(successTag, TupleTagList.of(deadLetterTag));

// Write dead letter elements to separate sink
outputTuple.get(deadLetterTag).apply(BigQuery.write(...));

// Process the successful element differently.
PCollection success = outputTuple.get(successTag);
```

```
final TupleTag successTag;
final TupleTag deadLetterTag;
PCollection input = /* ... */;

PCollectionTuple outputTuple = input.apply(ParDo.of(new DoFn(){
    @Override
    void processElement(ProcessContext ctxt) {
        try {
            c.output(process(c.element));
        } catch(MyException ex) {
            // Optional Logging at debug level
            c.sideOutPut(deadLetterTag, c.element);
        }
    }
})).writeOutPutTags(successTag, TupleTagList.of(deadLetterTag));

// Write dead letter elements to separate sink
outputTuple.get(deadLetterTag).apply(BigQuery.write(...));

// Process the successful element differently.
PCollection success = outputTuple.get(successTag);
```

```
final TupleTag successTag;
final TupleTag deadLetterTag;
PCollection input = /* ... */;

PCollectionTuple outputTuple = input.apply(ParDo.of(new DoFn(){
    @Override
    void processElement(ProcessContext ctxt) {
        try {
            c.output(process(c.element));
        } catch(MyException ex) {
            // Optional Logging at debug level
            c.sideOutPut(deadLetterTag, c.element);
        }
    }
})).writeOutPutTags(successTag, TupleTagList.of(deadLetterTag));

// Write dead letter elements to separate sink
outputTuple.get(deadLetterTag).apply(BigQuery.write(...));

// Process the successful element differently.
PCollection success = outputTuple.get(successTag);
```

Demo

Thank you!

Q & A